

CSC263 - Week 5, Lecture 1

Cristyn Howard

Monday, February 5, 2018

Q: A webserver needs to reference a large list of blacklisted websites. The entire list cannot be stored in main memory, because it is too large. How can we store it?

A: In a BLOOM FILTER

Bloom Filters

- probabilistic (i.e. approximate) dictionary that stores F_S : a summary/fingerprint of a dynamic set S

$BF[0, \dots, m-1]$: array of m bits, initialized to 0.

$\{h_1, \dots, h_t\}$: set of $t < m$ hash functions where $h_i : U \rightarrow \{0, \dots, m-1\}$

m, t : parameters defining # of bits & # of hash functions, respectively

- Operations:

$\Theta(n)$ INSERT(F_S, x) : for i from 1 to t , $BF[h_i(x)] = 1$;

- * element passed to all hash functions, produces a t -tuple of indices $\in \{0, \dots, m-1\}$ that form fingerprint of x
- * bits at all indices in fingerprint of x permanently set to 1

$\Theta(n)$ SEARCH(F_S, x) : for i from 1 to t , if $(BF[h_i(x)] == 0)$ return FALSE, else return TRUE

- * Returns: $\{\text{FALSE} \equiv x \text{ definitely not in } F_S\} || \{\text{TRUE} \equiv x \text{ **probably** in } F_S\}$
- * If any index in the hash of x is 0, then x has definitely not been placed in S .
- * However, if x is not in S , and the union of all indices in the hashes of all elements stored in S contains all indices in the hash of x , then SEARCH will return a **false positive**.

- Benefits: very space and time efficient way to store and search large sets.
- Limitations:
 - cannot delete elements from S
 - SEARCH may return *false positives* - may say that x is likely in S when it is not.
- Use when: space constraints exist, false positives are acceptable, and when deleting is not necessary.

False Positive Probability Analysis

Assume n elements inserted. Compute probability that $\text{SEARCH}(\text{BF}, x)$ returns a false positive, i.e. a positive when x has not been inserted.

Recall that hash functions have uniform, independent probability distributions.